

Test-Time Adaptation for LLMs via Hypernetwork-Generated LoRA Weights

Midterm Progress Report

Andrews George Varghese & Dhruv Kapur & Raj Maheshwari
{ageorgev, dkapur, rajmahes}@andrew.cmu.edu
10-423/623 Generative AI Course Project

December 11, 2025

1 Abstract

We investigate test-time adaptation of large language models through hypernetwork-generated LoRA weights that enable instance-level parameter specialization. Our approach extends recent work on task-level LoRA generation to produce unique adapter parameters for each input during inference, eliminating the need for separate per-task adapters. At the midterm checkpoint, we have successfully implemented: (1) a comprehensive evaluation framework supporting both chat and non-chat model formats across four diverse benchmarks (GSM8K, NLI, ARC-Easy and BoolQ) with flexibility to extend to more benchmarks, (2) baseline LoRA fine-tuning pipelines for three models spanning 70M to 600M parameters (Pythia 70M, Gemma3 270M Instruct and Qwen3 0.6B), and (3) an initial implementation of the hypernetwork that generates and injects LoRA weights into an LM. Preliminary baseline results demonstrate expected performance for the various model sizes and their dataset specific PEFT counterparts, providing clear targets for our hypernetwork approach. We outline remaining work including hypernetwork training, comparative evaluation across model scales, and analysis of generated adapter diversity.

25 1 Introduction

Parameter-efficient fine-tuning through Low-Rank Adaptation (LoRA) [Hu et al., 2022] has become the dominant paradigm for adapting large language models to specialized tasks, reducing trainable parameters by up to 10,000x while maintaining performance comparable to full fine-tuning. However, conventional LoRA approaches require training separate adapters for each task in advance, creating engineering overhead and preventing dynamic adaptation to individual inputs.

Recent work by Sakana AI [Charakorn et al., 2025] demonstrated that hypernetworks can generate task-

specific LoRA adapters from natural language descriptions, achieving 98% of specialized adapter performance without task-specific training. Building on this foundation, we investigate **instance-level adaptation**: using hypernetworks to generate unique LoRA parameters for each input during inference rather than uniform task-level adapters.

This midterm report documents our progress toward this goal, including implementation of evaluation infrastructure across three model scales (Pythia 70M, Gemma 270M-Instruction Tuned, Qwen3 0.6B) and four diverse benchmarks (GSM8K, NLI, ARC-Easy, BoolQ), baseline establishment, and initial hypernetwork development. We present our experimental setup, preliminary results, and detailed plans for completing the project.

2 Dataset, Task, and Model Selection

2.1 Benchmark Suite

To enable comprehensive evaluation while managing computational constraints, we focus on four diverse benchmarks representing different reasoning and knowledge requirements:

Mathematical Reasoning: GSM8K. GSM8K Cobbe et al. [2021] provides 8,500 grade-school math word problems requiring multi-step arithmetic reasoning (2 to 8 steps per solution). This tests the model’s ability to decompose problems and perform accurate numerical computations.

Natural Language Inference: SNLI. The Stanford Natural Language Inference (SNLI) dataset Bowman et al. [2015] contains 570,000 human-annotated sentence pairs labeled for entailment, contradiction, or neutral relationships. This requires the model to understand the 2 sentences and their relationship when seen as premise and hypothesis.

Science Reasoning: ARC-Easy. ARC-Easy Clark et al. [2018] contains 2,376 grade-school science ques-

74 tions that require reasoning beyond direct fact retrieval.
75 This benchmark tests scientific knowledge in a multi-
76 step reasoning setting.

77 **Question Answering: BoolQ.** The Boolean Ques-
78 tions Dataset (BoolQ) Clark et al. [2019] contains
79 around 16,000 naturally occurring yes/no questions
80 about Wikipedia passages, collected from real Google
81 search queries. Each question is paired with a para-
82 graph that contains the information needed to answer
83 it, thus testing if the model can perform reading com-
84 prehension and binary classification.

85 These four benchmarks provide diversity across rea-
86 soning types (mathematical, logical, scientific, ex-
87 tractive), output formats (free-form generation for
88 GSM8K, three-way classification for SNLI, multiple-
89 choice for ARC-Easy, yes/no for BoolQ) while remain-
90 ing tractable for our computational budget. More im-
91 portantly, it acts as an orthogonal set of datasets over
92 which we will train our hyper-network model. We aim
93 to cover as many such orthogonal ideas during training
94 while remaining within our computational budget, so
95 that the hyper-network is able to generalize to unseen
96 tasks.

97 2.2 Evaluation Metrics

98 We use task-appropriate metrics: match-with-tolerance
99 accuracy for GSM8K (requiring correct numerical
100 answers within some tolerance), and accuracy for
101 SNLI (three-way classification), ARC-Easy (multiple-
102 choice), and SQuAD (true/false).

103 2.3 Model Selection

Table 1: Model configurations for the three models evaluated in this work.

| Model | d | n_{layers} | h_q/h_{kv} |
|-----------------|------|---------------------|--------------|
| Pythia 70M | 512 | 6 | 8/8 |
| Gemma 3 270M-IT | 640 | 18 | 4/1 |
| Qwen3 0.6B | 1024 | 28 | 16/8 |

104 Rather than evaluating a single model, we adopt
105 a multi-model approach to understand how instance-
106 level adaptation scales across different model sizes and
107 training paradigms. Our three models span an order
108 of magnitude in parameter count and include both base
109 and instruction-tuned variants:

110 **Pythia 70M.** The smallest model in our suite, Pythia
111 70M [Biderman et al., 2023] is a decoder-only trans-
112 former trained on the Pile dataset. As a base pretrained
113 model without instruction tuning, it provides insight
114 into whether instance-level adaptation can effectively
115 specialize models with limited innate task-following

116 capabilities. Its small size enables rapid iteration dur-
117 ing hypernetwork development.

118 **Gemma 3 270M Instruct.** Gemma 3 270M-IT
119 [Gemma, 2025] represents the smallest, yet reason-
120 ably intelligent, instruction-tuned model we could find.
121 At approximately 4x the parameters of Pythia 70M, it
122 tests whether instance-level adaptation provides bene-
123 fits when the base model already possesses instruction-
124 following capabilities, but may not perform too well on
125 the benchmarks.

126 **Qwen3-0.6B.** Qwen3-0.6B is the largest model we
127 are considering, at approximately 9x the size of Pythia
128 70M. This model is extremely intelligent given its
129 small size, and it will be interesting to see if test-time
130 augmentation can further improve its performance.

131 This multi-model strategy enables us to assess: (1)
132 whether hypernetwork-generated adapters have a simi-
133 lar behavior on base and instruction-tuned models, (2)
134 how adaptation benefits scale with model capacity, and
135 (3) whether our approach generalizes across different
136 model architectures.

137 3 Related Work

138 Our work builds upon three interconnected research ar-
139 eas: parameter-efficient fine-tuning, hypernetwork ar-
140 chitectures, and meta-learning for fast adaptation.

141 **Low-Rank Adaptation.** LoRA [Hu et al., 2022] in-
142 troduced parameter-efficient fine-tuning by represent-
143 ing weight updates as $\Delta W = BA$ where $B \in \mathbb{R}^{d \times r}$
144 and $A \in \mathbb{R}^{r \times k}$ with rank $r \ll \min(d, k)$. This ap-
145 proach reduces trainable parameters by orders of mag-
146 nitude while matching full fine-tuning performance, es-
147 tablishing the foundation for modern efficient adapta-
148 tion methods.

149 **Hypernetworks for Parameter Generation.** The
150 paradigm of using neural networks to generate weights
151 for other networks was formalized by Ha et al. [2017],
152 demonstrating end-to-end training via backpropaga-
153 tion. HyperFormer Mahabadi et al. [2021] extended
154 this to transformers by generating adapter parameters
155 conditioned on task embeddings, layer positions, and
156 module types, achieving strong multi-task performance
157 with minimal per-task parameters (0.29%).

158 **Meta-Learning Foundations.** Model-Agnostic
159 Meta-Learning (MAML) [Finn et al., 2017] introduced
160 bi-level optimization for rapid task adaptation through
161 few gradient steps. While MAML requires gradi-
162 ent computation during adaptation, hypernetwork ap-
163 proaches generate parameters in a single forward pass,
164 representing a complementary trade-off between adap-
165 tation quality and computational efficiency.

166 **Task-Level LoRA Generation.** Most directly rel-
167 evant to our work, Text-to-LoRA [Charakorn et al.,

2025] demonstrated that hypernetworks can generate complete LoRA adapters from natural language task descriptions, compressing hundreds of task-specific adapters into a single model. HyperDecoders Iivison and Peters [2022] proposed treating parameter generation as autoregressive sequence generation conditioned on instructions, though focused on encoder-decoder architectures rather than decoder-only language models.

Our Contribution. While Text-to-LoRA operates at task-level granularity (one adapter per task description), we investigate *instance-level adaptation* that generates unique LoRA parameters for each input. This eliminates the need for task labels or descriptions while enabling fine-grained specialization to individual examples, potentially improving performance on heterogeneous data within tasks.

4 Methods

4.1 Baseline: Task-Specific LoRA

Our primary baseline is standard task-specific LoRA fine-tuning, applied independently to each benchmark dataset. For each task, we attach LoRA adapters of rank $r = 2$, and $\alpha = 8$ to all linear projection layers in the transformer. This configuration provides a highly parameter-efficient setting while remaining expressive enough to capture task-specific behavior. LoRA is applied to the query and value projection matrices except for Pythia, which has a single projection matrix for queries, keys and values.

We train using the default **AdamW** optimizer, a learning rate of 5×10^{-5} , per-device batch size of 2 with gradient accumulation of 2 (effective batch size 4 using 1 device), and a training budget of 500 gradient steps. These hyperparameters were chosen to ensure stability under low-rank adaptation, reduce overfitting on small datasets.

Overall, this per-task LoRA fine-tuning baseline represents a strong and competitive parameter-efficient adaptation method. It serves as the performance baseline for evaluating our hypernetwork-generated LoRA approach, which must be within acceptable range of these independently optimized adapters while training only once.

4.2 Hypernetwork Architecture

Our hypernetwork generates LoRA weight matrices conditioned on individual input sequences. The architecture consists of three components:

Input Encoding. We pass the input prompt through the frozen base LLM and extract the final layer’s last token representation as the semantic embedding, which

is the input to the hypernetwork. This self-encoding design offers two key advantages: First, it simplifies the implementation by requiring only a single tokenizer and leveraging the HuggingFace Transformers API without additional encoder dependencies. Second, and more importantly, it provides a strong inductive bias—the hypernetwork predicts LoRA weights based on the same internal representations that the target model will process, supposedly creating an alignment between the encoding space and the weight space.

Conditioning Information. Following Text-to-LoRA [Charakorn et al., 2025], we incorporate learned embeddings for layer position, module identity (for example q_proj , v_proj , etc.) and LoRA matrix identity (A or B). For a target layer ℓ , module m and matrix $k \in \{A, B\}$:

$$\mathbf{c}_{\ell,m,k} = \mathbf{h}_{\text{input}} + \mathbf{e}_{\text{layer}}^{(\ell)} + \mathbf{e}_{\text{module}}^{(m)} + \mathbf{e}_{\text{matrix}}^{(k)}$$

where $\mathbf{e}_{\text{layer}} \in \mathbb{R}^{n_{\text{layers}} \times d_h}$, $\mathbf{e}_{\text{module}} \in \mathbb{R}^{n_{\text{modules}} \times d_h}$ and $\mathbf{e}_{\text{matrix}} \in \mathbb{R}^{2 \times d_h}$ are learned embeddings, with d_h being the internal hypernetwork dimension, n_{layers} the number of transformer blocks in the target LM and n_{module} the cardinality of the set of target modules to apply LoRA to.

Weight Generation. The conditioned representation passes through a shared MLP Ψ . Each pair of target module-matrix has its own head (owing to the fact that different modules may have different matrix dimensions), with a total of $n_{\text{modules}} \times 2$ heads. Each head produces the number of outputs required for the corresponding matrix in the corresponding module.

$$\mathbf{z}_{\ell,m,k} = \Psi(\mathbf{c}_{\ell,m,k}) \quad (1)$$

$$A_{\ell,m} = W_{m,A} \mathbf{z}_{\ell,m,A} + b_{m,A} \quad (2)$$

$$B_{\ell,m} = W_{m,B} \mathbf{z}_{\ell,m,B} + b_{m,B} \quad (3)$$

Weight Injection: The target linear layers in the LM are replaced with a custom `DynamicLoraLinear` layer that allows dynamic injection of the predicted A and B weights as needed.

4.3 Training Strategy

We train the hypernetwork end-to-end by backpropagating through the frozen base language model. For each training example (x, y) :

1. Predict LoRA weights $\{B_{\ell,m}, A_{\ell,m}\}$ from input x
2. Inject generated adapter weights to the frozen base model
3. Compute language modeling loss $\mathcal{L}(y|x; \theta_{\text{base}}, \{B_{\ell,m} A_{\ell,m}\})$ and backpropagate gradients, updating only the hypernetwork

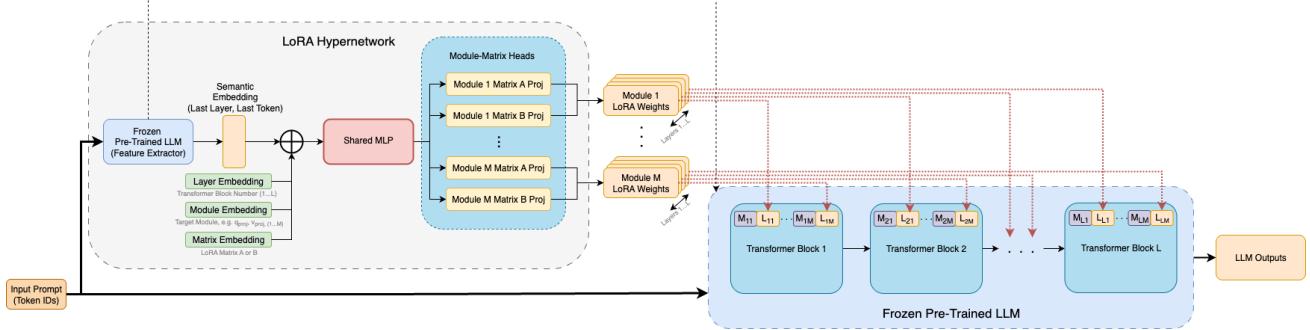


Figure 1: Overview of the TaskWeaver architecture. The hypernetwork (left) takes an input prompt and processes it through a frozen pre-trained LLM to extract a semantic embedding of the prompt (taken as the last layer representation of the last token). These embeddings are added to learned layer, module and matrix embeddings (A and B), then passed through a shared MLP backbone. Module-specific prediction heads generate LoRA weight matrices (A and B) for different modules (e.g., q_proj , v_proj , etc.) across all transformer layers. The generated LoRA weights are dynamically injected into the frozen base LLM (right) at inference time, adapting the model’s behavior on a per-input basis without requiring gradient-based optimization.

254 This approach directly optimizes the hypernetwork
 255 to generate adapters that maximize base model performance
 256 on training examples. The specific hyperparameters for training the hypernetwork are currently being
 257 worked on.

259 **Important Consideration** An important consideration
 260 comes from handling different LoRA weights
 261 for each batch element during training. Text-to-LoRA
 262 sidesteps this by using the same task description for all
 263 the elements in a batch. However, since our model con-
 264 ditions the LoRA adapters on the instance, instead of
 265 predicting $A \in \mathbf{R}^{r \times d}$ and $B \in \mathbf{R}^{d \times r}$, the hypernetwork
 266 predicts $A \in \mathbf{R}^{b \times r \times d}$ and $B \in \mathbf{R}^{b \times d \times r}$, which are ap-
 267 propriately multiplied with the inputs using BMM in
 268 DynamicLoraLinear.

5.2 Computational Resources

We have spread our training across multiple resources, including personal GPUs (an RTX 3070 and an M4 Pro Macbook), and 3090s and A100s rented from Vast.ai. Baseline LoRA training requires approximately 10-15 minutes on the A100 for the Qwen 0.6B per benchmark.

5.3 Experimental Results

Table 2 presents baseline performance across all four benchmarks and three model scales.

6 Plan Going Forward

We are currently on schedule with the original proposal plan and intend to complete the following for the final report.

6.1 Baseline Evaluation (Week +1)

Andrew, Raj: We plan to evaluate the baseline on more datasets, specifically ones used in Charakorn et al. [2025] such as HellaSwag [Zellers et al., 2019], Winogrande [Sakaguchi et al., 2020], Arc-C [Clark et al., 2018], etc. We will also evaluate the models on the all the benchmarks multiple times in order to get statistically significant results.

6.2 Hypernetwork Training and Evaluation (Week +1,2)

Dhruv, Raj: The collections of datasets will be combined into a superdataset in order to train the Hypernet-

269 5 Experimental Setup and Infra- 270 structure

271 5.1 Evaluation Framework

272 We implemented a flexible evaluation pipeline support-
 273 ing both chat and non-chat model formats. In particu-
 274 lar, it let’s us configure our tests using simple YAML
 275 files, including the test dataset and coverage, and paths
 276 to models and LoRA parameters if necessary. We also
 277 control temperature of generation. When the model is
 278 marked as a chat model, we modify its prompts accord-
 279 ingly (setting user, assistant and system roles, and uti-
 280 lizing tokenizer’s `apply_chat_template`).

Table 2: Baseline results comparing pretrained base models against task-specific LoRA fine-tuned models across three model scales and four benchmarks.

| Model | Variant | GSM8K | NLI | ARC-Easy | BoolQA |
|---------------|-------------------|--------------|------------|-----------------|---------------|
| Pythia 70M | Base | 0.38% | 6.30% | 10.69% | 13.64% |
| | + LoRA FT | 1.44% | 33.66% | 22.94% | 41.80% |
| | + Hypernet (Ours) | TBD | TBD | TBD | TBD |
| Gemma 270M-IT | Base | 8.50% | 31.87% | 0.8% | 41.81% |
| | + LoRA FT | 22.73% | 32.37% | 24.58% | 53.98% |
| | + Hypernet (Ours) | TBD | TBD | TBD | TBD |
| Qwen 0.6B | Base | 60.8% | 31.9% | 29.7% | 62.2% |
| | + LoRA FT | 72.78% | 75.64% | 76.56% | 63.24% |
| | + Hypernet (Ours) | TBD | TBD | TBD | TBD |

work. The specific architecture of the hypernetwork’s shared MLP is subject to change as we experiment with different architectures. Once these are finalized, we will do end-to-end training of the hypernetwork with our three models, and compare its performance with the baselines.

5 runs/model × 3 models × \$0.7/hr for just the small models we’re currently looking at)

This allocation would significantly strengthen our conclusions about the viability, generality, and scalability of instance-level LoRA adaptation across diverse tasks and model sizes.

6.3 Predicted Adapter Analysis (Week +2)

Andrew, Dhruv: Once the hypernetwork is trained and evaluated, we will analyze the impact of various factors on the generated adapters, such as input prompt, layer index, target module, etc.

8 Conclusion

At the midterm checkpoint, we have made substantial progress toward our goal of instance-level test-time adaptation via hypernetwork-generated LoRA weights. We have successfully implemented comprehensive evaluation infrastructure supporting both chat-formatted (Gemma, Qwen) and base (Pythia) models, established baseline performance targets across three model scales (70M–600M parameters) and four diverse benchmarks, and developed a hypernetwork architecture that dynamically injects LoRA weights into a base LM for instance-level adaptation. The specific hyperparameters and architecture choice is currently being optimized.

Our multi-model approach provides unique insights into how instance-level adaptation scales with model capacity and interacts with instruction tuning. The initial experiments set up strong baselines, providing clear targets for hypernetwork performance evaluation.

The remaining work focuses on completing hypernetwork training across all three models, comprehensive comparative evaluation across multiple datasets (the 4 currently reported and another 2-4 datasets), and analysis of generated adapter characteristics including cross-model behavior and instance-level diversity. We remain confident in achieving our core objectives while preparing contingency plans to ensure meaningful contributions regardless of final performance outcomes.

375 References

- 376 Stella Biderman, Hailey Schoelkopf, Quentin Gregory
377 Anthony, Herbie Bradley, Kyle O'Brien, Eric Hal-
378 lahan, Mohammad Aflah Khan, Shivanshu Purohit,
379 USVSN Sai Prashanth, Edward Raff, et al. Pythia:
380 A suite for analyzing large language models across
381 training and scaling. In *International Conference on*
382 *Machine Learning*, pages 2397–2430. PMLR, 2023.
- 383 Samuel R. Bowman, Gabor Angeli, Christopher Potts,
384 and Christopher D. Manning. A large annotated
385 corpus for learning natural language inference. In
386 Lluís Màrquez, Chris Callison-Burch, and Jian Su,
387 editors, *Proceedings of the 2015 Conference on*
388 *Empirical Methods in Natural Language Process-*
389 *ing*, pages 632–642, Lisbon, Portugal, September
390 2015. Association for Computational Linguistics.
391 doi: 10.18653/v1/D15-1075. URL <https://aclanthology.org/D15-1075>.
- 393 Rujikorn Charakorn, Edoardo Cetin, Yujin Tang,
394 and Robert Tjarko Lange. Text-to-LoRA: In-
395 stant transformer adaption. In *Forty-second*
396 *International Conference on Machine Learning*,
397 2025. URL <https://openreview.net/forum?id=zWskCdu3QA>.
- 399 Christopher Clark, Kenton Lee, Ming-Wei Chang, and
400 Tom Kwiatkowski. Boolq: Exploring the surprising
401 difficulty of natural yes/no questions. In *Proce-*
402 *edings of the 2019 Conference of the North American*
403 *Chapter of the Association for Computational Lin-*
404 *guistics*, 2019.
- 405 Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot,
406 Ashish Sabharwal, Carissa Schoenick, and Oyvind
407 Tafjord. Think you have solved question answering?
408 try arc, the ai2 reasoning challenge. In *Proce-*
409 *edings of the 2018 Conference on Empirical Methods*
410 *in Natural Language Processing*, 2018.
- 411 Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian,
412 Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias
413 Plappert, Jerry Tworek, Jacob Hilton, Reiichiro
414 Nakano, Christopher Hesse, and John Schulman.
415 Training verifiers to solve math word problems.
416 *arXiv preprint arXiv:2110.14168*, 2021.
- 417 Chelsea Finn, Pieter Abbeel, and Sergey Levine.
418 Model-agnostic meta-learning for fast adaptation of
419 deep networks. In *Proceedings of the 34th Inter-*
420 *national Conference on Machine Learning*, volume 70 of *Proceedings of Machine Learning Re-*
421 *search*, pages 1126–1135. PMLR, 2017.
- Team Gemma. Gemma 3. 2025. URL <https:// goo.gle/Gemma3Report>. 423
424
- David Ha, Andrew M Dai, and Quoc V Le. Hyper- 425
networks. In *International Conference on Learn- 426
ing Representations*, 2017. URL <https:// openreview.net/forum?id=rkpACellx>. 427
428
- Edward J Hu, Yelong Shen, Phillip Wallis, Zeyuan 429
Allen-Zhu, Yuanzhi Li, Shean Wang, Lu Wang, 430
and Weizhu Chen. LoRA: Low-rank adapta- 431
tion of large language models. In *Inter- 432
national Conference on Learning Representations*, 433
2022. URL <https:// openreview.net/forum?id=nZeVKeFYf9>. 434
435
- Hamish Ivison and Matthew E. Peters. Hyperde- 436
coders: Instance-specific decoders for multi-task 437
nlp, 2022. URL <https:// arxiv.org/abs/2203.08304>. 438
439
- Rabeeh Mahabadi, Sebastian Ruder, Mostafa De- 440
hghani, and James Henderson. Parameter-efficient 441
multi-task fine-tuning for transformers via shared 442
hypernetworks. In *Proceedings of the 59th Annual* 443
Meeting of the Association for Computational Lin- 444
guistics and the 11th International Joint Conference 445
on Natural Language Processing (Volume 1: Long 446
Papers), pages 565–576. Association for Compu- 447
tational Linguistics, 2021. 448
- Keisuke Sakaguchi, Ronan Le Bras, Chandra Bhaga- 449
vatula, and Yejin Choi. Winogrande: An adversarial 450
winograd schema challenge at scale. In *Proceedings* 451
of the AAAI Conference on Artificial Intelligence, 452
volume 34, pages 8732–8740, 2020. 453
- Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali 454
Farhadi, and Yejin Choi. Hellaswag: Can a machine 455
really finish your sentence? In *Proceedings of the* 456
57th Annual Meeting of the Association for Compu- 457
tational Linguistics, 2019. 458